

Class 08: Breast Cancer Mini Proj

Ryan Ellis (PID: A17673864)

Table of contents

Background	1
Data Import	1
Exploratory Data Analysis	3
Principle Component Analysis (PCA)	4
Lets visualize	5
Varaince	7
Communicating PCA Results	8
Hierarchical Clustering	9
Combine Methods	10
Prediction	15

Background

The goal of this mini-project is for you to explore a complete analysis using the unsupervised learning techniques covered in class.

We will analyze a data set describing a fine needle aspiration (FNA) of a breast mass.

Data Import

The data is made available as a CSV file for download. We can use the `read.csv()` function:

```
# IMPORTANT TO SET ROW NAMES EQUAL TO ONE
wisc.df<- read.csv("WisconsinCancer.csv", row.names=1)
```

```
head(wisc.df, 3)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean		
842302	M	17.99	10.38	122.8	1001		
842517	M	20.57	17.77	132.9	1326		
84300903	M	19.69	21.25	130.0	1203		
		smoothness_mean	compactness_mean	concavity_mean	concave.points_mean		
842302		0.11840	0.27760	0.3001	0.14710		
842517		0.08474	0.07864	0.0869	0.07017		
84300903		0.10960	0.15990	0.1974	0.12790		
		symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se	
842302		0.2419		0.07871	1.0950	0.9053	8.589
842517		0.1812		0.05667	0.5435	0.7339	3.398
84300903		0.2069		0.05999	0.7456	0.7869	4.585
		area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302		153.40	0.006399	0.04904	0.05373	0.01587	
842517		74.08	0.005225	0.01308	0.01860	0.01340	
84300903		94.03	0.006150	0.04006	0.03832	0.02058	
		symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302		0.03003		0.006193	25.38	17.33	
842517		0.01389		0.003532	24.99	23.41	
84300903		0.02250		0.004571	23.57	25.53	
		perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302		184.6	2019	0.1622	0.6656		
842517		158.8	1956	0.1238	0.1866		
84300903		152.5	1709	0.1444	0.4245		
		concavity_worst	concave.points_worst	symmetry_worst			
842302		0.7119		0.2654	0.4601		
842517		0.2416		0.1860	0.2750		
84300903		0.4504		0.2430	0.3613		
		fractal_dimension_worst					
842302		0.11890					
842517		0.08902					
84300903		0.08758					

We need to exclude the `diagnosis` column for analysis as this is the answer we are trying to achieve.

```
# This will give us everything but the first column
wisc.data <- wisc.df[,-1]
diagnosis<- as.factor(wisc.df$diagnosis)
```

Exploratory Data Analysis

Q1. How many observations are in this dataset?

```
nrow(wisc.df)
```

```
[1] 569
```

There are 569 patients or observations.

Q2. How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
diagnosis  
  B    M  
357 212
```

There are 357 benign and 212 are malignant.

Q3. How many variables/features in the data are suffixed with `_mean`?

We can use the `grep()` function:

```
colnames(wisc.data)
```

```
[1] "radius_mean"      "texture_mean"  
[3] "perimeter_mean"   "area_mean"  
[5] "smoothness_mean"  "compactness_mean"  
[7] "concavity_mean"    "concave.points_mean"  
[9] "symmetry_mean"     "fractal_dimension_mean"  
[11] "radius_se"         "texture_se"  
[13] "perimeter_se"      "area_se"  
[15] "smoothness_se"     "compactness_se"  
[17] "concavity_se"      "concave.points_se"  
[19] "symmetry_se"       "fractal_dimension_se"  
[21] "radius_worst"      "texture_worst"  
[23] "perimeter_worst"   "area_worst"  
[25] "smoothness_worst"  "compactness_worst"  
[27] "concavity_worst"   "concave.points_worst"  
[29] "symmetry_worst"    "fractal_dimension_worst"
```

```
length(grep("_mean", colnames(wisc.data)))
```

```
[1] 10
```

There are 10 variables with `_mean`

Principle Component Analysis (PCA)

Need to scale data before PCA. `scale=TRUE` in the argument of `prcomp()`.

```
## In general will always want to include scaling, will make the plots easier to interpret as
wisc.pr<- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

Around 44.27% of the data is captured via PC1

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?


To describe 70% of the data you need at least the first three PC's so PC1-PC3

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

To get to 90% cumulative proportion it is necessary to have 7 principal components (PC1-PC7)

Lets visualize


```
biplot(wisc.pr)
```



Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

The sheer amount of data that this plot is trying to visualize. It is extremely difficult to visualize much of anything within the plot due to the number of points and different things occurring.


```
library(ggplot2)
ggplot(wisc.pr$x)+
  aes(PC1, PC2, col=diagnosis)+
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

That both PC1/PC2 and PC1/PC3 are separated via diagnosis on PC1 axis, but both heavily overlap in PC2 and PC3. PC1 captures the most amount of variance.

```
ggplot(wisc.pr$x)+
  aes(PC1, PC3, col=diagnosis)+
  geom_point()
```




Varaince

We will make a scree plot, which will show us the variance that each PC indicates. Basically creates a exponential graph (hill), that will show when adding new PC's doesn't really make sense anymore.

```
pr.var<-(wisc.pr$sdev^2)
```

```
pve<- ((pr.var)/(30))
```

```
plot(c(1,pve), xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



Communicating PCA Results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

```
wisc.pr$rotation[,1]
```

radius_mean	texture_mean	perimeter_mean
-0.21890244	-0.10372458	-0.22753729
area_mean	smoothness_mean	compactness_mean
-0.22099499	-0.14258969	-0.23928535
concavity_mean	concave.points_mean	symmetry_mean
-0.25840048	-0.26085376	-0.13816696
fractal_dimension_mean	radius_se	texture_se
-0.06436335	-0.20597878	-0.01742803
perimeter_se	area_se	smoothness_se
-0.21132592	-0.20286964	-0.01453145
compactness_se	concavity_se	concave.points_se
-0.17039345	-0.15358979	-0.18341740


symmetry_se	fractal_dimension_se	radius_worst
-0.04249842	-0.10256832	-0.22799663
texture_worst	perimeter_worst	area_worst
-0.10446933	-0.23663968	-0.22487053
smoothness_worst	compactness_worst	concavity_worst
-0.12795256	-0.21009588	-0.22876753
concave.points_worst	symmetry_worst	fractal_dimension_worst
-0.25088597	-0.12290456	-0.13178394

The proportional component of PC1 that is `concave.points_mean` is `-.261`, and there is no other component with as much magnitude. This means that this feature has the biggest influence on where the point (patient) is placed in the PC1 axis.

Hierarchical Clustering

```
wisc.hclust <- hclust(dist(scale(wisc.data)))
plot(wisc.hclust)
```


Cluster Dendrogram



```
dist(scale(wisc.data))
hclust (*, "complete")
```

Combine Methods

```
d<-dist(wisc.pr$x[,1:4])
wisc.pr.hclust<- hclust(d, method="ward.D2")
plot(wisc.pr.hclust)
abline(h=75, col="green")
```



```
grps<- cutree(wisc.pr.hclust, h=75)
table(grps)
```

```
grps
  1  2
171 398
```

```
# These results are similar to the diagnosis, where we cluster the two groups into Malignant
```

How does this clustering in `grps()` correspond to the expert **diagnosis**

```
table(diagnosis)
```

```
diagnosis
  B   M
357 212
```


```
table(diagnosis, grps)
```

```
      grps
diagnosis 1   2
  B      6 351
  M     165 47
```

Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h=19, col="red", lty=2)
```

Cluster Dendrogram




```
dist(scale(wisc.data))
hclust (*, "complete")
```

The height at which the normal clustering model has 4 clusters without PCA is around 19.

Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

```
method.hclust<- hclust(dist(scale(wisc.data)), method="average")  
plot(method.hclust)
```


Cluster Dendrogram



```
dist(scale(wisc.data))  
hclust (*, "average")
```

```
method.hclust<- hclust(dist(scale(wisc.data)), method="complete")  
plot(method.hclust)
```


Cluster Dendrogram



```
dist(scale(wisc.data))  
hclust (*, "complete")
```


```
method.hclust<- hclust(dist(scale(wisc.data)), method="single")  
plot(method.hclust)
```

Cluster Dendrogram



```
dist(scale(wisc.data))  
hclust (*, "single")
```

```
method.hclust<- hclust(dist(scale(wisc.data)), method="ward.D2")
plot(method.hclust)
```



```
dist(scale(wisc.data))
hclust (*, "ward.D2")
```

The best method is **ward.D2**, also average is alright. They both aim to spread the data put so it is easier to interpret the clusters. but **ward.D2** does a much better job of this when looking at the above plots of the different methods.

Q13. How well does the newly created hclust model with two clusters separate out the two “M” and “B” diagnoses?

It does decently when we compare it to the expert diagnosis, it shows that we have 6 false positive for benign and 47 false negatives for malignant. This is shown in the table above plotted grps versus diagnosis.

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.hclust.clusters` and `wisc.pr.hclust.clusters`) with the vector containing the actual diagnoses.

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=2)
table(diagnosis,wisc.hclust.clusters)
```

```

wisc.hclust.clusters
diagnosis  1  2
B 357    0
M 210    2

```

Before using PCA, we can see that there is almost no separation between the two types, benign and malignant. Although when using PCA we have some false negative and positives, without it we can't really even interpret anything.

Prediction

We will now utilize the `predict()` function to combine our previous PCA model and some new data (cancer cell data)

```

url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc

```

```


      PC1      PC2      PC3      PC4      PC5      PC6      PC7
[1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
      PC8      PC9      PC10     PC11     PC12     PC13     PC14
[1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
[2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
      PC15     PC16     PC17     PC18     PC19     PC20
[1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,] 0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
      PC21     PC22     PC23     PC24     PC25     PC26
[1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
[2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
      PC27     PC28     PC29     PC30
[1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
[2,] -0.001134152 0.09638361 0.002795349 -0.019015820

```

```

plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")

```



Q16. Which of these new patients should we prioritize for follow up based on your results?

When viewing the plot, and based off the previous conclusions drawn from the PCA plots, a malignant patient(tumor), is likely to be in the negative portion on the x axis (PC1). Thus when comparing patient 1 and 2 because patient 2 lies on the negative side amongst the malignant patients I would predict this patient to be prioritized.